

# Reinforcement Learning Assignment-1

## Logic For Environment

- Using gym as abstract class for my environment
- action\_space = discrete values including [0,0],[0,1],[1,0],[1,1],[0,-1],[-1,0],[-1,1],[-1,-1]
- Observation\_space = all pair of discrete value of positions in state + all pair of velocities (total = 6048)
- reset() :-
  - choose a random start state
- Collision Detector :- I am assuming that the car is moving in a straight line connecting its initial position and position after translating with v velocity. If the object hits the wall or is outside the track and not finished, we reset. If it finish, we will stop the episode.

## On Policy Algorithm Used

On-policy first-visit MC control (for  $\varepsilon$ -soft policies), estimates  $\pi \approx \pi_*$

Algorithm parameter: small  $\varepsilon > 0$

Initialize:

$\pi \leftarrow$  an arbitrary  $\varepsilon$ -soft policy

$Q(s, a) \in \mathbb{R}$  (arbitrarily), for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$  empty list, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following  $\pi$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair  $S_t, A_t$  appears in  $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$ :

Append  $G$  to  $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \operatorname{argmax}_a Q(S_t, a)$

(with ties broken arbitrarily)

For all  $a \in \mathcal{A}(S_t)$ :


$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

- I initialized the q as np.ones() of dimension (number of states, number of actions)
- I initialized the pi as np.ones()/number\_of\_actions of dimension (number of states, number of actions)
- I am choosing epsilon as 0.1 for my algorithm and gamma as 0.1
- number\_of\_episodes = 1,00,000 (time taken:- 2.5 hrs average)
- for each episode
  - returns,state\_indexes,action\_indexes = run episode
  - returns is the sum G without discount
  - for each timestep
    - $q[\text{state}, \text{action}] = q[\text{state}, \text{action}] + \text{gamma} * (\text{return at that timestep} - q[\text{state}, \text{action}])$


- $\text{greedyActions} = \text{np.argmax of } q \text{ for each state}$
- set value of  $\pi$  at all state, greedy action as  $1 - \epsilon + \epsilon / \text{number\_of\_actions}$
- set all other as  $\epsilon / \text{number\_of\_actions}$
- Then set policy value of all action for each state as 1 if that action in  $\text{argmax}$  of state

### On policy plots:-


- start state 1




- start state 2



- start state 3



- Action using policy for in each state



### Off Policy Algorithm Used

### Off-policy MC control, for estimating $\pi \approx \pi_*$

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$Q(s, a) \leftarrow \text{arbitrary}$

$C(s, a) \leftarrow 0$

$\pi(s) \leftarrow \operatorname{argmax}_a Q(s, a)$  (with ties broken consistently)

Repeat forever:

$b \leftarrow \text{any soft policy}$

Generate an episode using  $b$ :

$S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T, S_T$

$G \leftarrow 0$

$W \leftarrow 1$

For  $t = T - 1, T - 2, \dots$  down to 0:

$G \leftarrow \gamma G + R_{t+1}$

$C(S_t, A_t) \leftarrow C(S_t, A_t) + W$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)} [G - Q(S_t, A_t)]$

$\pi(S_t) \leftarrow \operatorname{argmax}_a Q(S_t, a)$  (with ties broken consistently)


If  $A_t \neq \pi(S_t)$  then exit For loop

$W \leftarrow W \frac{1}{b(A_t|S_t)}$


- I initialized the q as np.ones() of dimension (number of states, number of actions)
- I initialized the pi as np.ones()/number\_of\_actions of dimension (number of states, number of actions)
- C = zeros for dimension (number of states,number of action)
- soft\_policy = result of on policy
- num\_itations = 1,00,000
- gamma = 0.1
- for each iteration
  - w = 1
  - states,returns,actions = run episode
  - for each timestep
    - c(state,action) = 1 + w
    - q(state,action) += gamma\* (g - q[state,action])/(W/c[state,action])
    - set policy[state, greedyAction] as 1 rest as 0. Here greedy action is np.argmax(q[state])
    - if action != greedy action, move to next episode
    - else W = W/(soft\_policy[state,action])

### Off Policy Plots

- starting state 1




- starting state 2






- starting state 3



### Convergence Test


I am checking the difference in policy before the iteration and after the iteration and then plotting the graph between episode number and difference.

For On Policy Method:-



As we can see difference is getting very close to 0 when episode count increase, so we can say on policy is converging

For Off-policy



As number of iterations increase by 40000, difference become close to 0. So we can say, off policy is also converging.

## References

- [https://github.com/BY571/Medium\\_Code\\_Examples/blob/master/Gridworld/Monte%20Carlo%20Methods%20Examples.ipynb](https://github.com/BY571/Medium_Code_Examples/blob/master/Gridworld/Monte%20Carlo%20Methods%20Examples.ipynb)
- <https://medium.com/analytics-vidhya/monte-carlo-methods-in-reinforcement-learning-part-1-on-policy-methods-1f004d59686a>
- <https://github.com/vojtamolda/reinforcement-learning-an-introduction>
- <https://towardsdatascience.com/solving-racetrack-in-reinforcement-learning-using-monte-carlo-control-bdee2aa4f04e>